

Localized Data - How to Map It With Hibernate

Let's first take a look at the table model before we discuss how you can map it with JPA and Hibernate. To make that easier to understand, I will use the following example:

We want to create a marketplace in which suppliers can offer their products. The marketplace supports the languages German and English. The supplier can provide the name and description of a product in both languages.

As so often, you can model this in various ways. Shantanu Kher created a great overview of different options and discussed their advantages and disadvantages on the [vertabelo blog](#).

Even though the popularity of these approaches varies, I have seen all of them in real life. In my experience, the most commonly used ones are:

1. Using separate columns for each language in the same database table, e.g., modeling the columns `description_en` and `description_de` to store different translations of a product description.
2. Storing translated fields in a separate table. That would move the `description_en` and `description_de` columns to a different table. Let's call it `LocalizedProduct`.

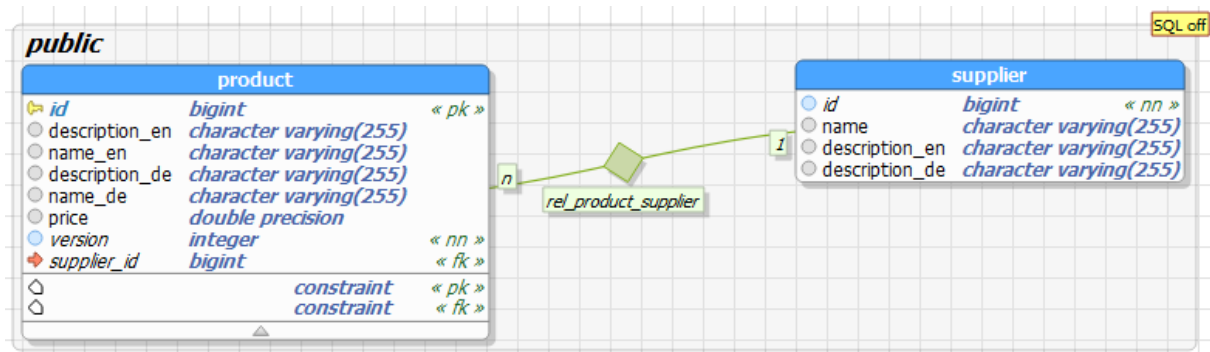
Let's take a closer look at both options.

Separate Language Columns in Each Table

The general idea of this approach is simple. For each localized attribute and language you need to support, you add an extra column to your table. Depending on the number of supported languages and localized attributes, this can result in a vast amount of additional columns. If you want to translate 4 attributes into 5 different languages, you would need to model $4 \times 5 = 20$ database columns.

In the previously described example, you need 4 database columns to localize the product name and description. You use the columns `description_en` and `description_de` to persist the different translations of the product description. The columns `name_en` and `name_de` to store the localized product name.

Localized Data - How to Map It With Hibernate



Creating Your Entity Mappings

As you have seen in the previous diagram, using separate columns for each translation results in a straightforward table model. The same is true for the entity mapping.

The `id` attribute is of type `Long` and maps the primary key. The `@GeneratedValue` annotation tells Hibernate to use a database sequence to generate unique primary key values. In this example, I use Hibernate's default sequence. But as I showed in a previous article, you can easily provide your own sequence.

The `version` attribute is used for optimistic locking and provides a highly scalable way to avoid concurrent updates. I explain it in more details in my Hibernate Performance Tuning Online Training.

The `supplier_id` attribute defines the owning side of a many-to-one association to the `Supplier` entity. As for all to-one associations, you should make sure to set the FetchType to LAZY to avoid unnecessary queries and performance problems.

The `nameDe`, `nameEn`, `descriptionDe`, and `descriptionEn` attributes just map each of the localized columns. That might result in a lot of attributes, but it is also a simple and efficient way to handle localized data.

Localized Data - How to Map It With Hibernate

```
@Entity
public class Product {

    @Id
    @GeneratedValue(strategy = GenerationType.AUTO)
    private Long id;

    @Version
    private int version;

    @ManyToOne(fetch = FetchType.LAZY)
    private Supplier supplier;

    private Double price;

    @Column(name = "name_de")
    private String nameDe;

    @Column(name = "name_en")
    private String nameEn;

    @Column(name = "description_de")
    private String descriptionDe;

    @Column(name = "description_en")
    private String descriptionEn;

    ...
}
```

Localized Data - How to Map It With Hibernate

Pros & Cons of Entities with Separate Language Columns

As you have seen, adding a dedicated column for each translation to your table:

- Is very easy to implement in the table model,
- Is very easy to map to an entity and
- Enables you to fetch all translations with a simple query that doesn't require any JOIN clauses.

But on the downside:

- this mapping might require a lot of database columns if you need to translate multiple attributes into various languages,
- fetching an entity loads translations that you might not use in your use case and
- you need to update the database schema if you need to support a new language.

In my experience, the inflexibility of this approach is the biggest downside. If your application is successful, your users and sales team will request additional translations. The required schema update makes supporting a new language much harder than it should be. You not only need to implement and test that change, but you also need to update your database without interrupting your live system.

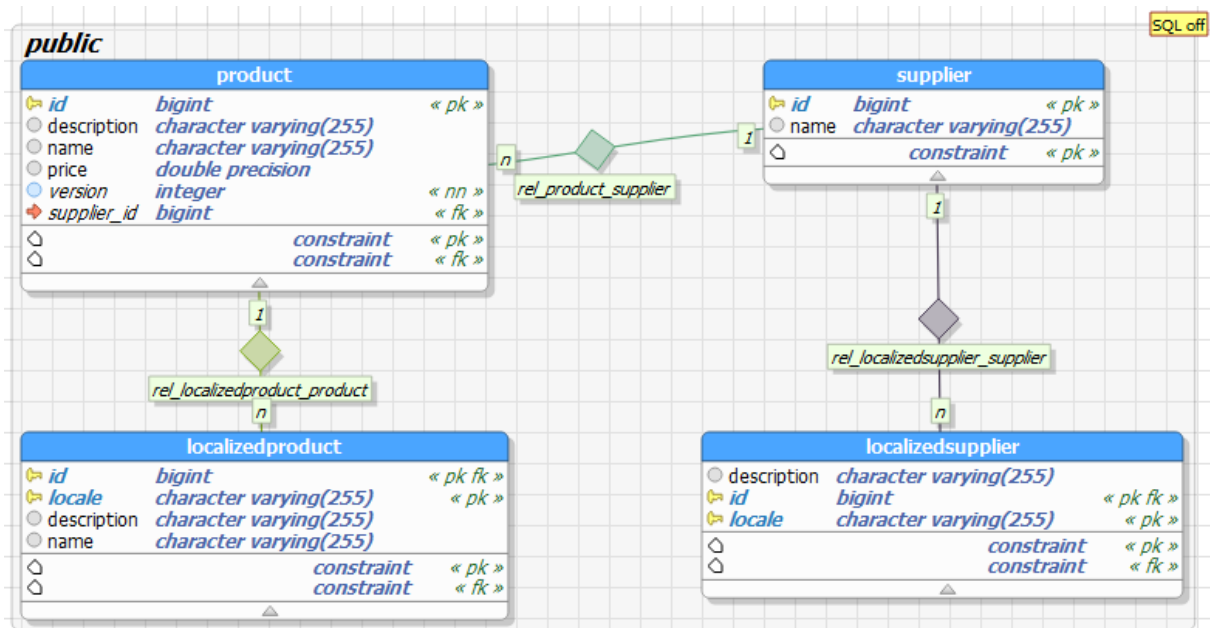
The next approach avoids these problems, and I, therefore, recommend it for most applications.

Different Tables and Entities for Translated and Non-Translated Fields

Instead of storing all the information in the same database table, you can also separate the translated and non-translated fields into 2 tables. That enables you to model a one-to-many association between the non-translated fields and the different localizations.

Here you can see a table model that applies this approach to the previously discussed example.

Localized Data - How to Map It With Hibernate



The LocalizedProduct table stores the different translations of the product name and description. As you can see in the diagram, that table contains a record for each localization of a product. So, if you want to store an English and a German name and description of your product, the LocalizedProduct table contains 2 records for that product. And if you're going to support an additional language, you only need to add another record to the LocalizedProduct table instead of changing your table model.

Creating Your Entity Mappings

The entity model is almost identical to the table model. You map the non-translated columns of the Product table to the Product entity and the translated columns of the LocalizedProduct table to the LocalizedProduct entity. And between these 2 entity classes, you can model a managed many-to-one association.

Entity with Translated Fields - The LocalizedProduct entity

The following mapping of the LocalizedProduct entity consists of a few mandatory and an optional part. Let's first talk about the mandatory mapping of the primary key and the association to the Product entity.

Localized Data - How to Map It With Hibernate

```
@Entity
@Cache(usage = CacheConcurrencyStrategy.TRANSACTIONAL)
public class LocalizedProduct {

    @EmbeddedId
    private LocalizedId localizedId;

    @ManyToOne
    @MapsId("id")
    @JoinColumn(name = "id")
    private Product product;

    private String name;

    private String description;

    ...
}
```

The LocalizedProduct entity represents the to-many side of the association. The Product product attribute, therefore, owns the relationship definition. The @JoinColumn annotation tells Hibernate to use the id column of the LocalizedProduct table as the foreign key column. And the @MapsId annotation defines that the primary key value of the associated Product entity is part of the composite primary key of the LocalizedProduct entity. It gets mapped to the id attribute of the primary key class.

As I explain in great details in the [Advanced Hibernate Online Training](#), you can map a composite primary key in various ways with JPA and Hibernate. In this example, I use an embedded id and an embeddable called LocalizedId.

Localized Data - How to Map It With Hibernate

As you can see in the following code snippet, the `LocalizedId` class is a basic Java class which implements the `Serializable` interface and is annotated with `@Embeddable`. And because you want to use it as an embedded id, you also need to make sure to implement the `equals` and `hashCode` methods.

```
@Embeddable
public class LocalizedId implements Serializable {

    private static final long serialVersionUID = 1089196571270403924L;

    private Long id;

    private String locale;

    public LocalizedId() {
    }

    public LocalizedId(String locale) {
        this.locale = locale;
    }

    ...
}
```

If you want to take it one step further, you might also want to cache the `LocalizedProduct` entity. You can do that by activating the cache in your `persistence.xml` configuration and by annotating the `LocalizedProduct` entity with JPA's `@Cacheable` or Hibernate's `@Cache` annotation. As I explain in my [Hibernate Performance Tuning Online Training](#), caching is a two-edged sword. It can provide substantial performance benefits but also introduce an overhead which can slow down your application. You need to make sure that you only change data that gets often read but only rarely changed. In

Localized Data - How to Map It With Hibernate

most applications, that's the case for the localized Strings. That makes them excellent candidates for caching.

Entity with Non-Translated Fields - The Product entity

After we mapped the LocalizedProduct table, which represents the different translations of the localized fields, it's time to work on the mapping of the Product table.

The only difference to the previous example is the mapping of the localized attributes. Instead of mapping an attribute for each translation, I'm using the localizations attribute. It maps the referencing side of the many-to-one association to the LocalizedProduct entity to a `java.util.Map`. This is one of the more advanced association mappings defined by the JPA specification, and I explained in great details in [How to map an association as a `java.util.Map`](#).

In this example, I use the locale attribute of the LocalizedProduct entity as the key and the LocalizedProduct entity as the value of the Map. The locale is mapped by the LocalizedId embeddable, and I need to specify the path `localizedId.locale` in the `@MapKey` annotation.

The mapping to a `java.util.Map` makes accessing a specific translation in your business code more comfortable. And it doesn't affect how Hibernate fetches the association from the database. In your JPQL or Criteria Queries, you can use this association in the same way as any other managed relationship.

Localized Data - How to Map It With Hibernate

```
@Entity
public class Product {

    @Id
    @GeneratedValue(strategy = GenerationType.SEQUENCE)
    private Long id;

    @Version
    private int version;

    @ManyToOne(fetch = FetchType.LAZY)
    private Supplier supplier;

    private Double price;

    @OneToMany(mappedBy = "product", cascade = {CascadeType.DETACH,
        CascadeType.MERGE, CascadeType.PERSIST, CascadeType.REFRESH},
        orphanRemoval = true)
    @MapKey(name = "localizedId.locale")
    @Cache(usage = CacheConcurrencyStrategy.TRANSACTIONAL)
    private Map<String, LocalizedProduct> localizations = new HashMap<>();

    ...

    public String getName(String locale) {
        return localizations.get(locale).getName();
    }

    public String getDescription(String locale) {
        return localizations.get(locale).getDescription();
    }
}
```

Localized Data - How to Map It With Hibernate

If you want to make your entity model more comfortable to use, you could activate `orphanRemoval` for the association. That is a general best practice for one-to-many associations that model a parent-child relationship in which the child can't exist without its parent. It tells your JPA implementation, e.g., Hibernate, to delete the child entity as soon as its association to the parent entity gets removed. I use it in this example to remove a `LocalizedProduct` entity as soon as it's no longer associated with a `Product` entity.

Another thing you could do to improve the usability of your entities is to provide getter methods that return the product name and description for a given locale. If you implement additional getter methods to return a localized name and description, you need to keep in mind that they are accessing a lazily fetched one-to-many association. That triggers an additional SQL statement if the association isn't already fetched from the database. You can avoid that by using a JOIN FETCH clause or an entity graph to initialize the association while loading your `Product` entity.

And if you activated the 2nd level cache on the `LocalizedProduct` entity, you should also annotate the localizations attribute with Hibernate's `@Cache` annotation. That tells Hibernate to cache the association between these 2 entities. If you miss this annotation, Hibernate will execute a query to retrieve the associated `LocalizedProduct` entities even though they might be already in the cache. That is another example of how complex caching with Hibernate can be. It's also one of the reasons why the Hibernate Performance Tuning Online Training includes a very detailed lecture about it.

Pros & Cons of Different Entities for Translated and Non-Translated Fields

Storing your translations in a separate table is a little more complicated, but it provides several benefits:

- Each new translation is stored as a new record in the `LocalizedProduct` table. That enables you to store new translations without changing your table model.

Localized Data - How to Map It With Hibernate

- Hibernate's 2nd level cache provides an easy way to cache the different localizations. In my experience, other attributes of an entity, e.g., the price, change more often than the translations of a name or description. It can, therefore, be a good idea to separate the localizations from the rest of the data to be able to cache them efficiently.

But the mapping also has a few disadvantages:

- If you want to access the localized attributes, Hibernate needs to execute an additional query to fetch the associated `LocalizedProduct` entities. You can avoid that by initializing the association when loading the `Product` entity.
- Fetching associated `LocalizedProduct` entities might load translations that you don't need for your use case.

Conclusion

Using additional columns to store the translations of a field, might seem like the most natural and obvious choice. But as I showed you in this article, it's very inflexible. Supporting an additional language requires you to change your table and your domain model.

You should, therefore, avoid this approach and store the translated and non-translated information in 2 separate database tables. You can then map each table to an entity and model a one-to-many association between them.

This approach allows you to add new translations without changing your domain and table model. But the mapping is also a little more complicated, and Hibernate needs to execute an additional query to retrieve the different localizations. You can avoid these queries by activating the 2nd level cache.